

```
>hadoop@sel-a1-216-25:~$ hadoop version
```

```
>hadoop@sel-a1-216-25:~$ ls
```

```
>hadoop@sel-a1-216-25:~$ cd hadoop-3.4.1
```

```
>hadoop@sel-a1-216-25:~/hadoop-3.4.1$ ls
```

```
bin include libexec      licenses-binary logs      NOTICE.txt sbin  
etc lib  LICENSE-binary LICENSE.txt  NOTICE-binary README.txt share
```

```
>hadoop@sel-a1-216-25:~/hadoop-3.4.1$ cd bin
```

```
>hadoop@sel-a1-216-25:~/hadoop-3.4.1/bin$ ls
```

```
container-executor hdfs  mapred.cmd      yarn  
hadoop            hdfs.cmd oom-listener    yarn.cmd  
hadoop.cmd        mapred  test-container-executor
```

```
>hadoop@sel-a1-216-25:~/hadoop-3.4.1/bin$ start-dfs.sh
```

```
Starting namenodes on [localhost]
```

```
Starting datanodes
```

```
Starting secondary namenodes [sel-a1-216-25]
```

```
>hadoop@sel-a1-216-25:~/hadoop-3.4.1/bin$ start-yarn.sh
```

```
Starting resourcemanager
```

```
Starting nodemanagers
```

```
>hadoop@sel-a1-216-25:~/hadoop-3.4.1/bin$ jps
```

```
10824 ResourceManager
```

```
10458 DataNode
```

```
10938 NodeManager
```

```
10315 NameNode
```

```
11132 Jps
```

```
10620 SecondaryNameNode
```

Code

```
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

import java.io.IOException;
import java.util.StringTokenizer;

public class WordCount {

    public static class TokenizerMapper extends Mapper<Object, Text, Text, IntWritable> {
        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context) throws IOException,
        InterruptedException {
            StringTokenizer itr = new StringTokenizer(value.toString());
            while (itr.hasMoreTokens()) {
                word.set(itr.nextToken().replaceAll("[^a-zA-Z]", "").toLowerCase()); // Clean words
                if (!word.toString().isEmpty()) {
                    context.write(word, one);
                }
            }
        }
    }

    public static class IntSumReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
        private IntWritable result = new IntWritable();

        public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException,
        InterruptedException {
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            result.set(sum);
            context.write(key, result);
        }
    }

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
```

```
    Job job = Job.getInstance(conf, "word count");
    job.setJarByClass(WordCount.class);
    job.setMapperClass(TokenizerMapper.class);
    job.setCombinerClass(IntSumReducer.class);
    job.setReducerClass(IntSumReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(IntWritable.class);
    FileInputFormat.addInputPath(job, new Path(args[0]));
    FileOutputFormat.setOutputPath(job, new Path(args[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}
```

Input.txt:

Anurag
Ganesh
Pranav
Ganesh
Anurag
Anurag

VSCODE Terminal

```
>hadoop@sel-a1-216-25:~/31452$ javac -classpath $(hadoop classpath) -d . WordCount.java
```

```
>hadoop@sel-a1-216-25:~/31452$ jar cf wc.jar WordCount*.class
```

```
>hadoop@sel-a1-216-25:~/31452$ hdfs dfs -put input.txt /user/input/
```

```
>hadoop@sel-a1-216-25:~/31452$ hadoop jar wc.jar WordCount input.txt output
```

```
>hadoop@sel-a1-216-25:~/31452$ hdfs dfs -cat output/part-r-00000
```

Output:

```
hdoop@sel-a1-216-25:~/31452$ hdfs dfs -cat output1/part-r-00000
```

```
anurag 3
```

```
ganesh 2
```

```
pranav 1
```